

CS118 Discussion 1B, Week 1

Boyan Ding

Outline

- Logistics
- Question on latency calculation
- Intro to network programming

TA

- Boyan Ding, PhD in Computer Networking
- Office hours:
 - Mon/Wed 9:30-10:30 am by Boyan Ding
 - Tuesday 1:30-3:30 pm by Zhehui Zhang
- Emails: dboyan@cs.ucla.edu
 - Please use [CS118] in subject or may be flagged as spam
 - Please include your name, UID in email.

Logistics

- Submit your signed Academic Integrity Agreement
- Grading breakdown
 - Homework: 20%
 - Programming projects: 20%
 - 3 quizzes: 60% (20% for each quiz)
 - Bonus credit: up to 15% for extra work in the programming projects (counting 3% in the final accumulated score)
- no late turn-in will be accepted for credit
- no make-up exams

Logistics: Homework

- Online submission to Gradescope only (course entry code: KYD66B). DEMO
 - Fill in your **UCLA ID** so that we can submit your score to CCLE
- Submission guidelines:
 - 1. **Hard deadline** on submission, so submit early! You can **resubmit** multiple times before the deadline, but the system will not accept submissions after the deadline.
 - 2. Each homework problem will have a dedicated **answering box** immediately below. Do **NOT** write your answers outside the box. Any answer outside the dedicated area may not get graded.
 - 3. You are encouraged to work out the problem on the PDF file directly **without altering the page layout in any way.**
 - 4. If you prefer handwriting or have to draw diagrams, you may scan the paper copy (e.g., using a smartphone app), convert it to a PDF file and then upload. It is **your** responsibility to upload a high-quality copy in black and white. Inaccessible answers will get low scores.

Logistics: Project

- Two projects (in C/C++):
 - A simple web server — get familiar with network programming;
 - Reliable data transfer — implement a simple user-level TCP-like transport protocol
- Test environment:
 - Ubuntu virtual machine

What are we learning in this course

Lectures:

Part 1: Introduction (2 lectures, text: Chapter 1)

Part 2: Application Layer (2 lectures, text: Ch.2)

-- introduction to socket programming is provided on Friday recitations

Part 3: Transport Layer (5 lectures, text Ch. 3)

**Project 1: April 23, Friday*

*Quiz 1: Cover Part 1, Part 2, and portion of Part 3 (up to sec. 3.4 (with reliable data transfer included))

Part 4: Network Layer (4 lectures, text: Ch. 4 and 5)

*Quiz 2: Cover portion of Part 3 (start from TCP), & Part 4

Part 6: Link Layer, LANs (4 lectures, text: Ch. 6)


Part 7: Wireless and Mobile Networks (1.5 lectures, text: Ch. 7)

Part 8: Network Security (1.5 lecture: Ch. 8)

*Project 2: June 4, Friday

*Quiz 3: Cover Parts 6, 7 & 8 (final exam slot: 6/11/2021).

Packet delay review



$$d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$$

d_{proc} : nodal processing

- check bit errors
- determine output link

d_{trans} : transmission delay:

- L : packet length (bits)
- R : link *bandwidth* (bps)
- $d_{\text{trans}} = L/R$

d_{queue} : queueing delay [[DEMO](#)]

- time waiting at output link for transmission
- depends on congestion level of router

d_{prop} : propagation delay:

- d : length of physical link
- s : propagation speed ($\sim 2 \times 10^8$ m/sec)
- $d_{\text{prop}} = d/s$

Question

- a. How long does it take a packet of length 1000 bytes to propagate over a link of distance 2500km, propagation speed 2.5×10^8 m/s, and transmission rate 2 Mbps?
- b. More generally, how long does it take a packet of length L to propagate over a link of distance d , propagation speed s , and transmission rate R bps?
- c. Does this delay depend on packet length?
- d. Does this delay depend on transmission rate?

Solution

- a. How long does it take a packet of length 1000 bytes to propagate over a link of distance 2500km, propagation speed 2.5×10^8 m/s, and transmission rate 2 Mbps?
- Ans: $(2500 \times 10^3) / (2.5 \times 10^8) = 0.01 \text{ s} = 10 \text{ ms}$
- b. More generally, how long does it take a packet of length L to propagate over a link of distance d , propagation speed s , and transmission rate R bps? Ans: d/s
- c. Does this delay depend on packet length? Ans: No
- d. Does this delay depend on transmission rate? Ans: No

Question

- Suppose Host A wants to send a large file to Host B. The path from Host A to Host B has three links, of rates $R_1=500\text{kbps}$, $R_2=2\text{Mbps}$, and $R_3=1\text{Mbps}$.
 - a. Assuming no other traffic in the network, what is the throughput for the file transfer?
 - b. Suppose the file is 4 million bytes. Dividing the file size by the throughput, roughly how long will it take to transfer the file to Host B?
 - c. Repeat (a) and (b), but now with R_2 reduce to 100kbps .

Solution

- a. Assuming no other traffic in the network, what is the throughput for the file transfer? **Ans: 500 kbps**
- b. Suppose the file is 4 million bytes. Dividing the file size by the throughput, roughly how long will it take to transfer the file to Host B? **Ans: $(4 \times 10^6) \times 8 / (500 \times 10^3) = 64$ seconds**
- c. Repeat (a) and (b), but now with R2 reduce to 100kpbs.
Ans: 320 seconds


Question cont'd

- Suppose Host A wants to send a large file to Host B. The path from Host A to Host B has three links, of rates $R_1=500\text{kbps}$, $R_2=100\text{kbps}$, and $R_3=1\text{Mbps}$.
 - d. Suppose the file is 200 bytes and it is segmented into two 100 bytes packets. What is the queuing delay for the second packet at host A, the first node, the second node?

Question

- Consider an HTTP client that wants to retrieve a Web document at a given URL. The IP address of the HTTP server is initially unknown. What transport and application layer protocols besides HTTP are needed in this scenario?

Network Programming



Network programming

- **What is the model for network programming?**
- Where are we programming?
- Which APIs can we use? How to use them?

Client-server model

- Asymmetric communication
 - Client — requests data:
 - Initiates communication
 - Waits for server's response
 - Server (Daemon) — responds data requests:
 - Discoverable by clients (e.g. IP address + port)
 - Waits for clients connection
 - Processes requests, sends replies

Demo: telnet

```
~$ telnet google.com 80
Trying 216.58.217.206...
Connected to google.com.
Escape character is '^]'.
GET / HTTP/1.1

HTTP/1.1 200 OK
Date: Fri, 12 Jan 2018 21:44:31 GMT
Expires: -1
Cache-Control: private, max-age=0
Content-Type: text/html; charset=ISO-8859-1
P3P: CP="This is not a P3P policy! See g.co/p3phelp for more info."
Server: gws
X-XSS-Protection: 1; mode=block
X-Frame-Options: SAMEORIGIN
Set-Cookie: 1P_JAR=2018-01-12-21; expires=Sun, 11-Feb-2018 21:44:31 GMT; path=/;
domain=.google.com
Set-Cookie: NID=12...J; expires=Sat, 14-Jul-2018 21:44:31 GMT; path=/; domain=.google.com; HttpOnly
Accept-Ranges: none
Vary: Accept-Encoding
Transfer-Encoding: chunked

754a
<!doctype html><html itemscope="" itemtype="http://schema.org/WebPage" lang="en"><head><meta
content="Search the world's information, including webpages, images, videos and more. Google has
many special features to help you find exactly what you're looking for." name="description">
```

Client-server model

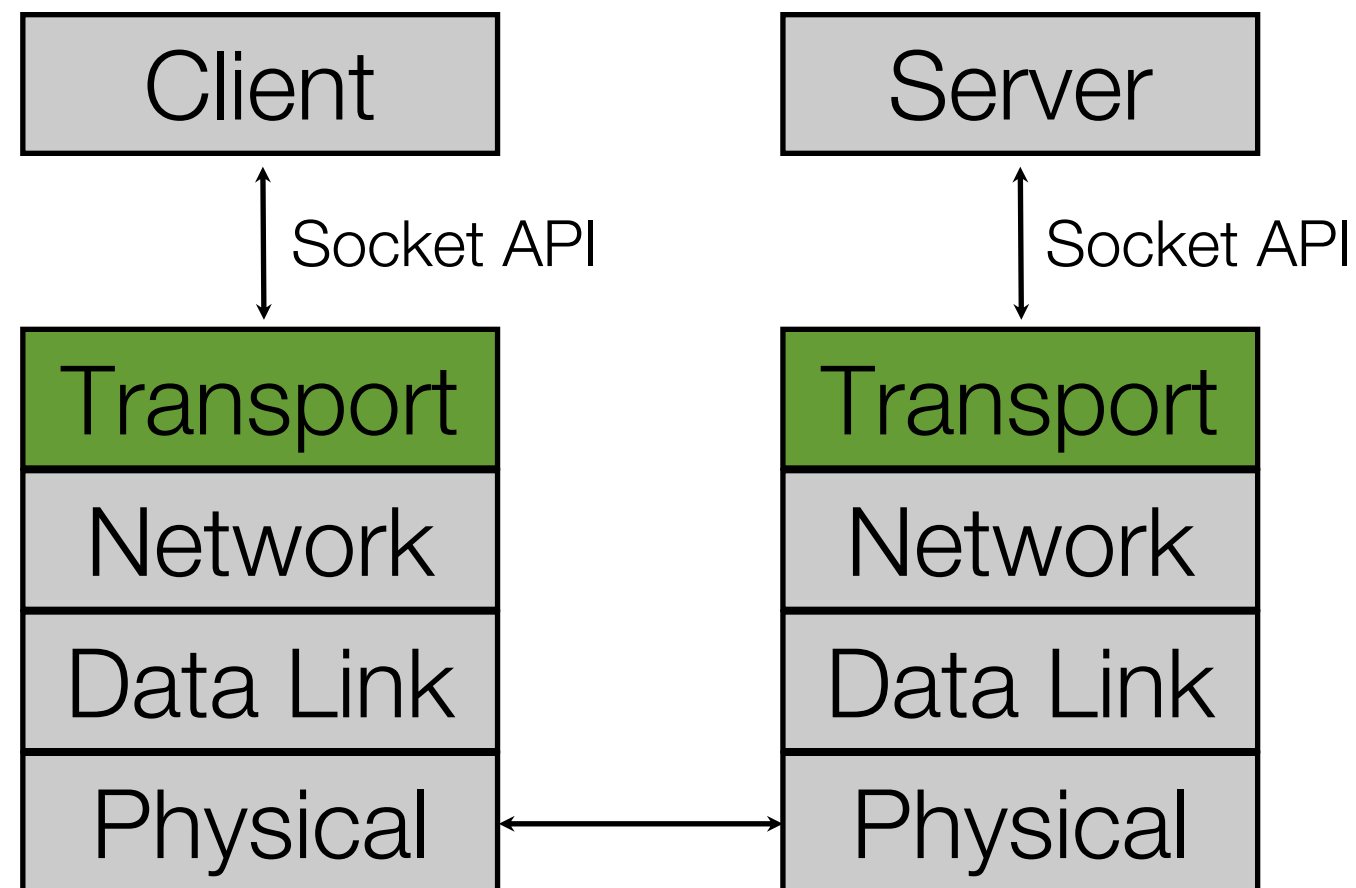
- Client and server are not disjoint
 - A client can be a server of another client
 - A server can be a client of another server
 - Example?
- Server's service model
 - Concurrent: server processes multiple clients' requests simultaneously
 - Sequential: server processes clients' requests one by one
 - Hybrid: server maintains multiple connections, but responses sequentially

Network programming

- What is the model for network programming?
- **Where are we programming?**
- Which APIs can we use? How to use them?

Which layer are we at?

- “Clients” and “servers” are programs at application layer
- Transport layer is responsible for providing communication services for application layer
- Basic transport layer protocols:
 - TCP
 - UDP



TCP: Transmission Control Protocol

- A connection is set up between client and server
- Reliable data transfer
 - Guarantee deliveries of all data
 - No duplicate data would be delivered to application
- Ordered data transfer
 - If A sends data D1 followed by D2 to B, B will also receive D1 before D2
- Data transmission: full-duplex byte stream (in two directions simultaneously)
- Regulated data flow: flow control and congestion control

UDP: User Data Protocol


- Basic data transmission service
 - Unit of data transfer: datagram (in variable length)
- No reliability guarantee
- No ordered delivery guarantee
- No flow control / congestion control

Network programming

- What is the model for network programming?
- Where are we programming?
- **Which APIs can we use? How to use them?**

Our secret weapon: socket programming APIs

- From Wikipedia: “A network socket is an endpoint of an inter-process communication flow across a computer network”
- A socket is a tuple of `<ip_addr:port>`
- Socket programming APIs help build the communication tunnel between applications and transport/network service
- We use TCP socket in this project



Socket: port number

- Port numbers are allocated and assigned by the IANA (Internet Assigned Numbers Authority)
- See RFC 1700 or <https://www.ietf.org/rfc/rfc1700.txt>

1-512	<ul style="list-style-type: none">• standard services (see <code>/etc/services</code>)• super-user only
513-1023	<ul style="list-style-type: none">• registered and controlled, also used for identity verification• super-user only
1024-49151	<ul style="list-style-type: none">• registered services/ephemeral ports
49152-65535	<ul style="list-style-type: none">• private/ephemeral ports

TCP socket: basic steps

- Create service
- Establish a TCP connection
- Send and receive data
- Close the TCP connection

TCP socket: service setup

TCP Client

TCP Server



TCP socket: service setup

TCP Client

TCP Server

socket()

TCP socket: service setup

TCP Client

TCP Server

socket()

bind()

TCP socket: service setup

TCP Client

TCP Server

socket()

bind()

listen()

TCP socket: service setup

TCP Client

TCP Server

socket()

bind()

listen()

accept()

TCP socket: service setup

TCP Client

TCP Server

socket()

bind()

listen()

accept()

blocked until
connection
from client

TCP socket: service setup

TCP Client

socket()

TCP Server

socket()

bind()


listen()

accept()


blocked until
connection
from client

TCP socket: establish connection

TCP Client




TCP Server




blocked until
connection
from client

TCP socket: send and receive data

TCP Client




TCP Server




blocked until
connection
from client

TCP socket: send and receive data

TCP Client



TCP Server




data (request)


blocked until
connection
from client

TCP socket: send and receive data

TCP Client



TCP Server

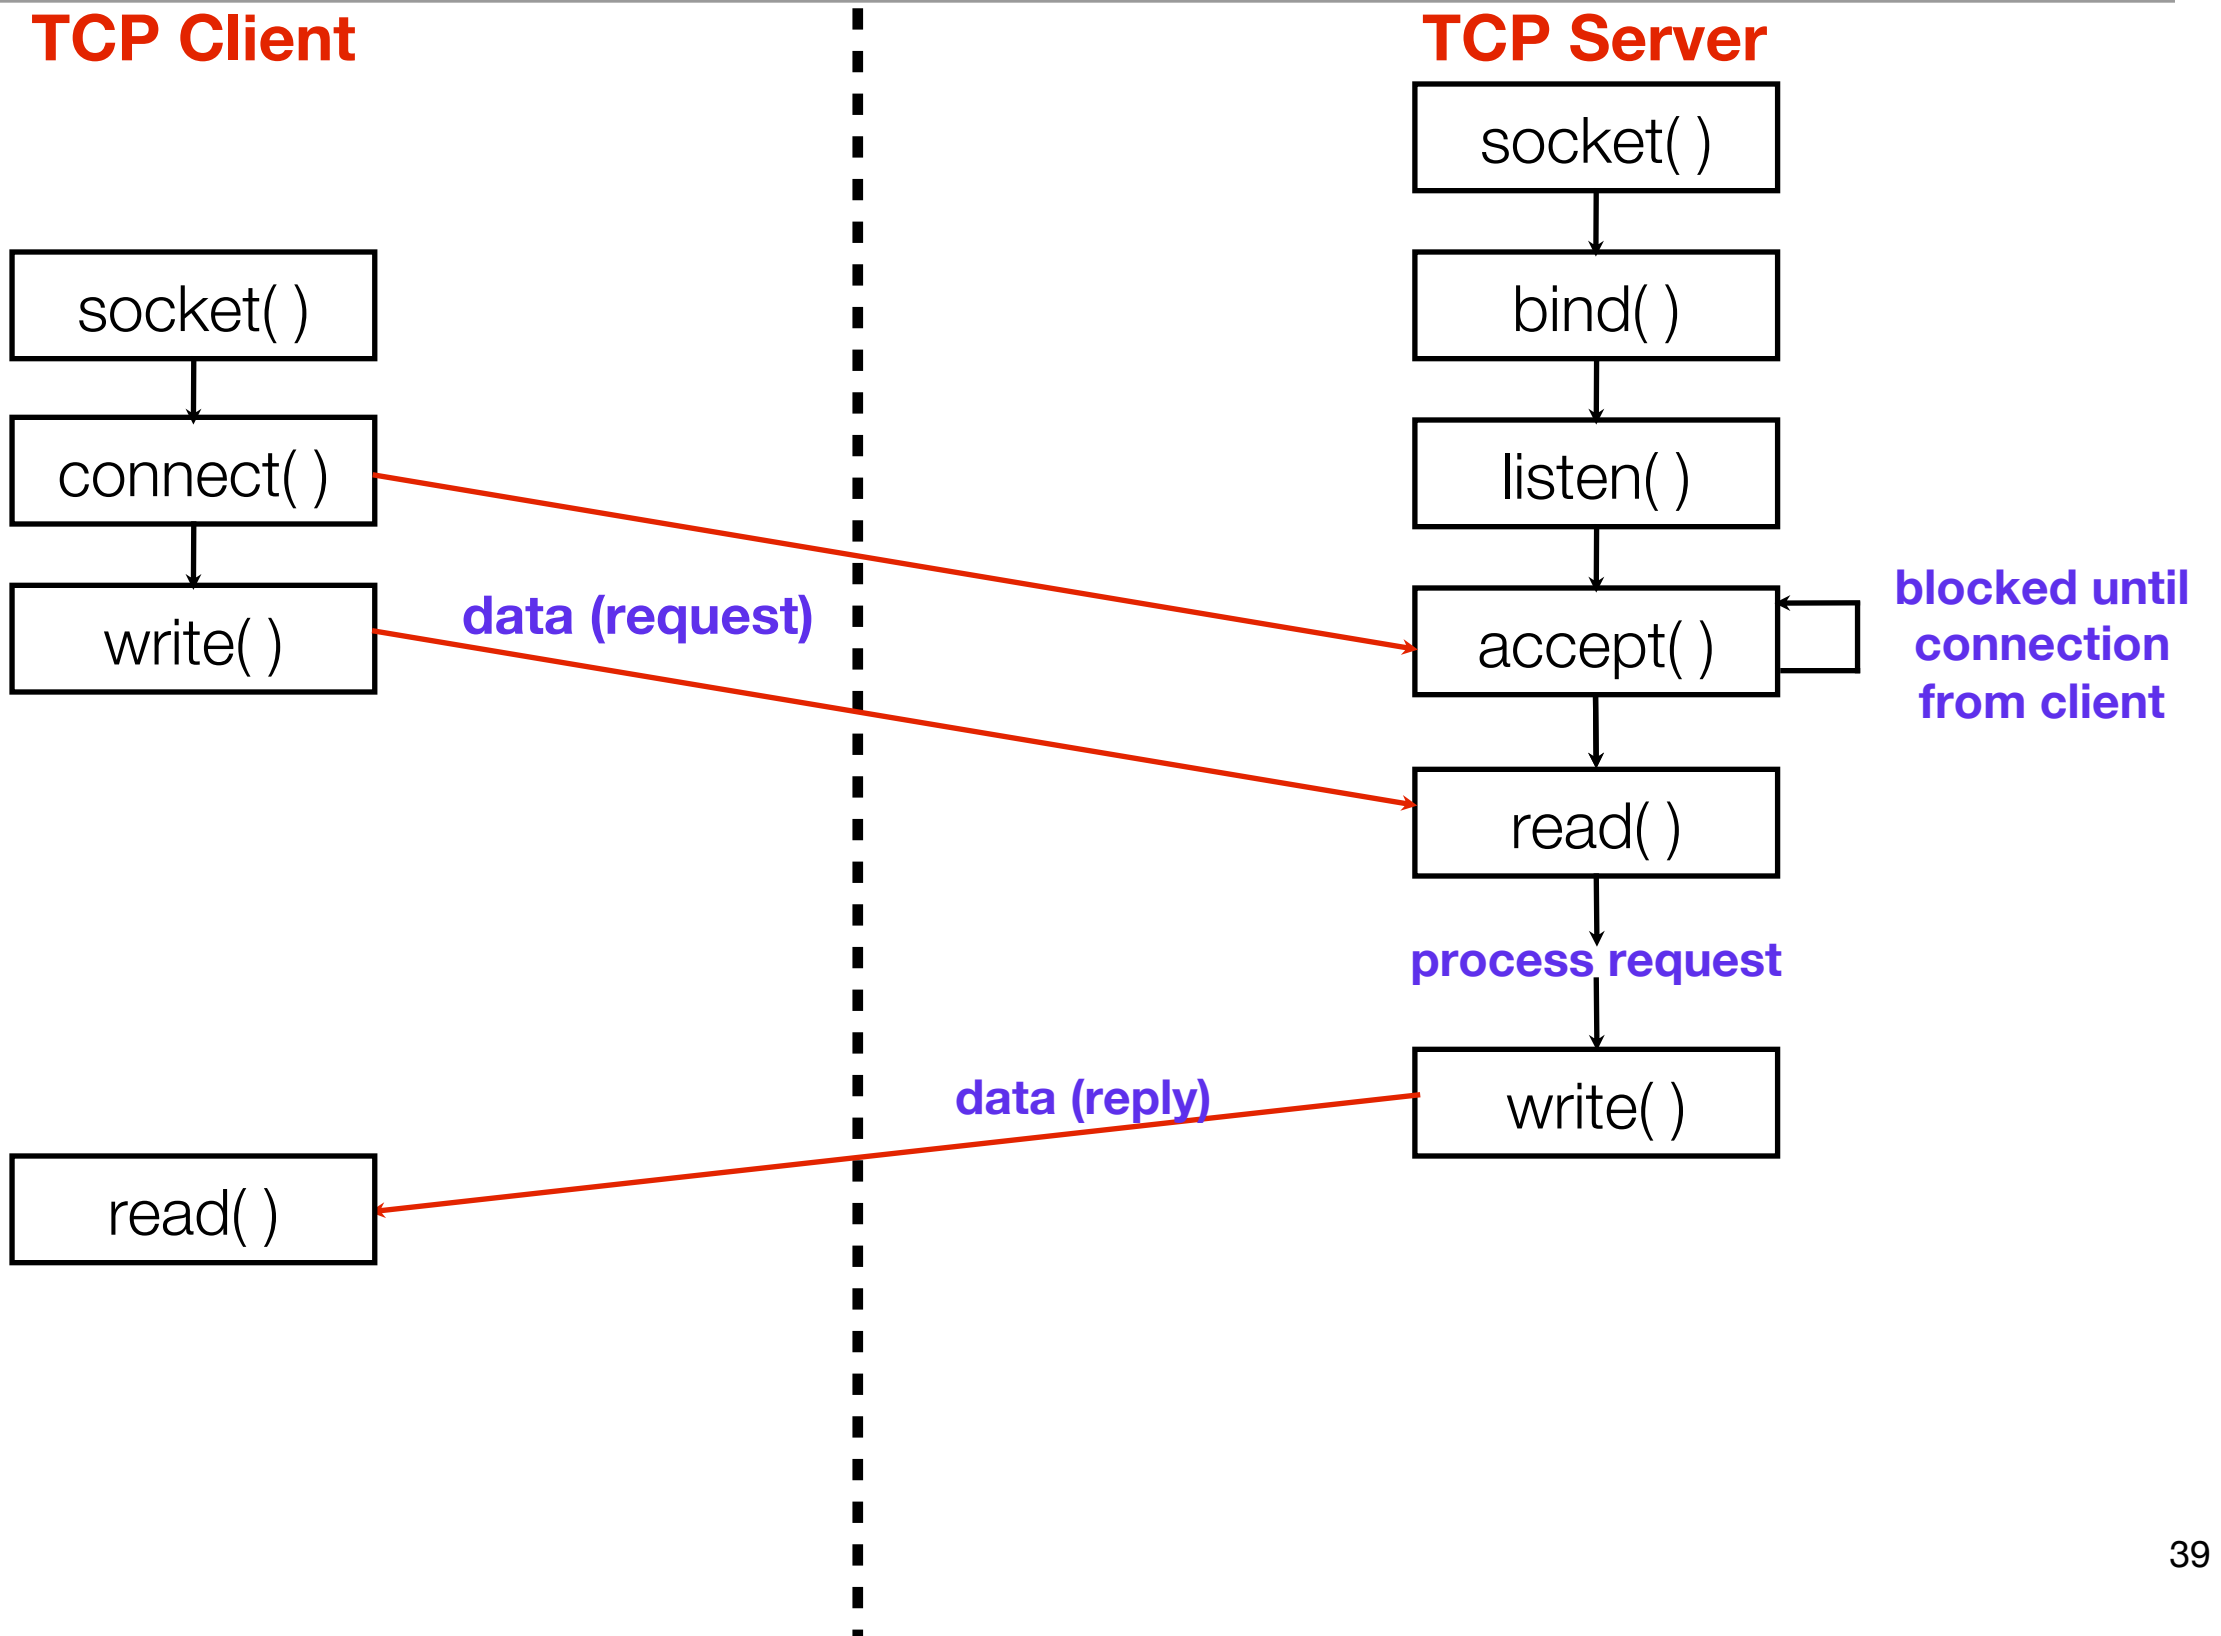


data (request)

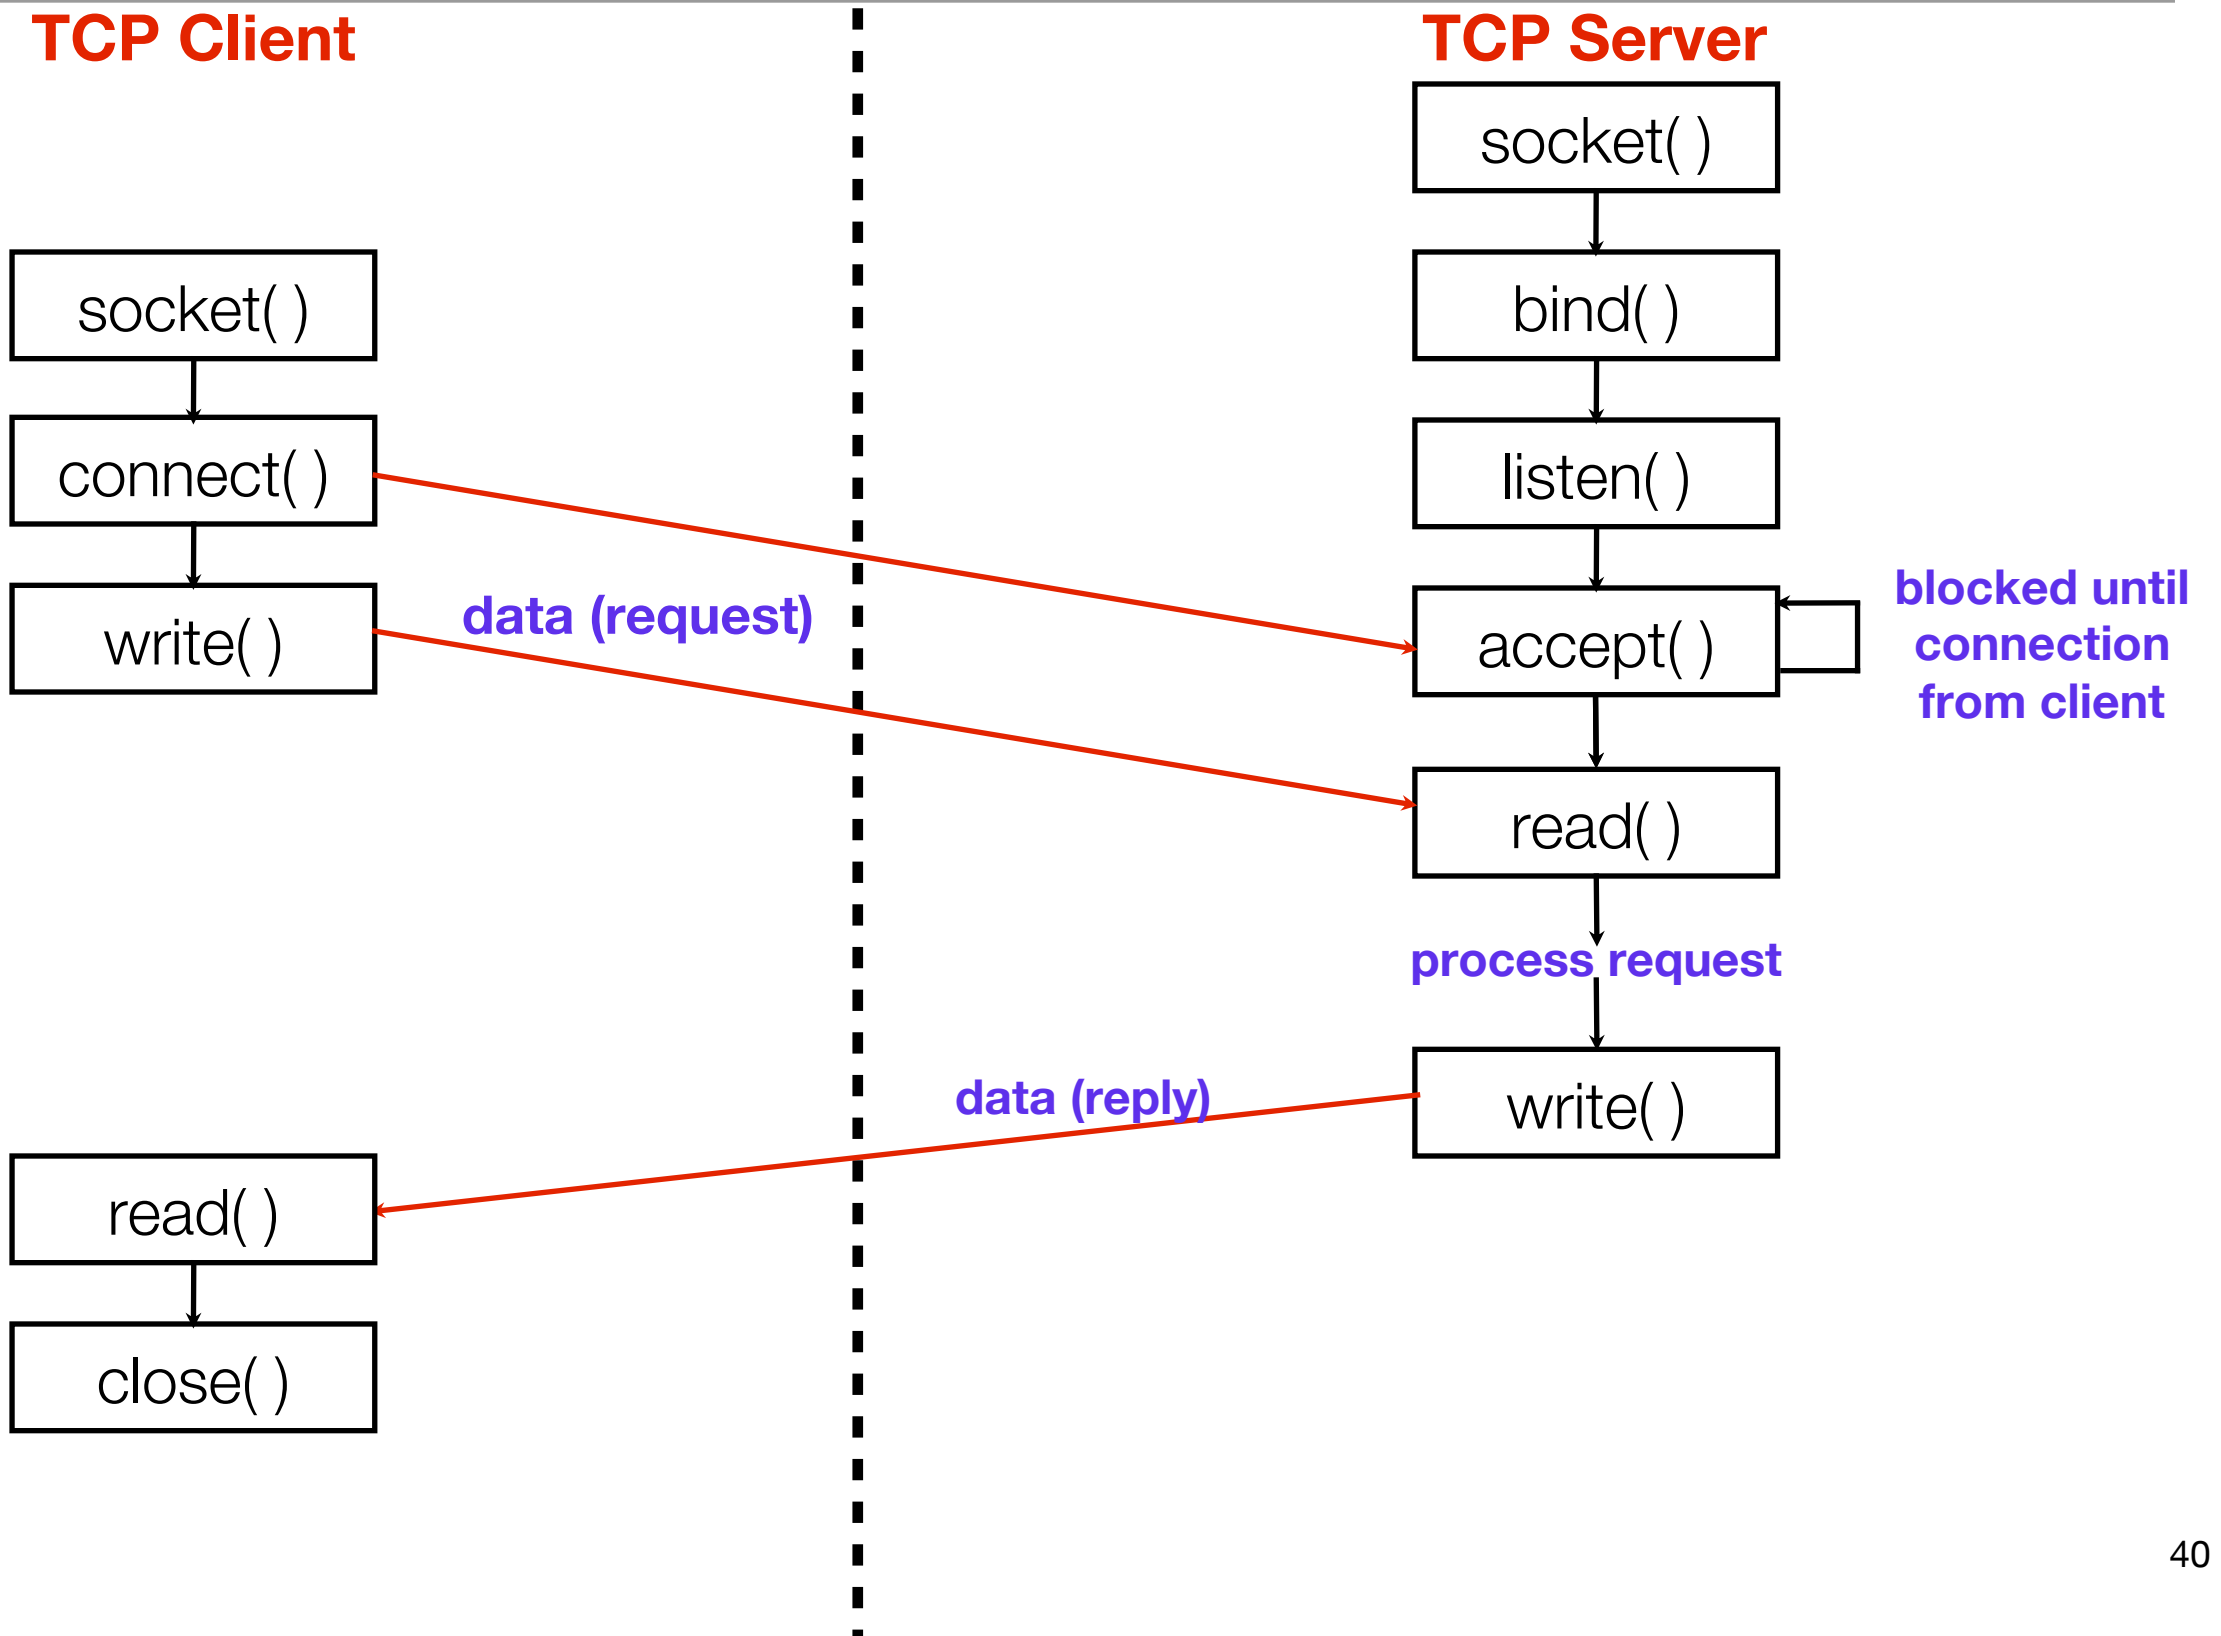
blocked until
connection
from client

process request

TCP socket: send and receive data



TCP socket: close connection



TCP socket: close connection

TCP Client

socket()

connect()

write()

read()

close()

TCP Server

socket()

bind()

listen()

accept()

read()

process request

write()

read()

close()

data (request)

data (reply)

blocked until
connection
from client

Socket programming API: syscalls

- **int socket(int domain, int type, int protocol);**
 - Create a socket
 - returns the socket descriptor or -1 (failure). Also sets errno upon failure
- **domain:** protocol family
 - **PF_INET** for IPv4, **PF_INET6** for IPv6, **PF_UNIX** or **PF_LOCAL** for Unix socket, **PF_ROUTE** for routing (You may also see AF_... version in some places, they basically mean the same thing)
- **type:** communication style
 - **SOCK_STREAM** for TCP (with **PF_INET**)
 - **SOCK_DGRAM** for UDP (with **PF_INET**)
- **protocol:** protocol within family, which is typically set to 0

Socket programming API: essential structs

- sockfd — socket descriptor. Just a regular `int`.
- sockaddr — socket address info
- sockaddr_in — yet another struct for the ‘internet’

```
struct sockaddr {
    unsigned short sa_family; // addr family, AF_xxx
    char sa_data[14]; // 14 bytes of proto addr
};
struct sockaddr_in { // used for IPv4 only
    short sin_family; // addr family, AF_INET
    unsigned short sin_port; // port number
    struct in_addr sin_addr; // internet address
    unsigned char sin_zero[8]; // zeros, same size as sockaddr
};
struct in_addr { // used for IPv4 only
    uint32_t s_addr; // 32-bit IPv4 address
};
```

Socket programming API: syscalls

- `int bind(int sockfd, struct sockaddr* myaddr, int addrlen);`

- Bind a socket to a local IP address and port number

- returns 0 on success, -1 and sets errno on failure

- **sockfd**: socket file descriptor returned by `socket ()`

- **myaddr**: includes IP address and port number

- **NOTE**: `sockaddr` and `sockaddr_in` are of same size, use `sockaddr_in` and convert it to `socketaddr`

- **sin_family**: protocol family, e.g. `AF_INET`

- **sin_port**: port number assigned by caller

- **sin_addr**: IP address

- **sin_zero**: used for keeping same size as `sockaddr`

- **addrlen**: `sizeof(struct sockaddr_in)`

```
struct sockaddr {
    short sa_family;
    char sa_data[14];
};

struct sockaddr_in {
    short sin_family;
    ushort sin_port;
    struct in_addr sin_addr;
    unsigned char sin_zero[8];
};
```

a pointer to a struct `sockaddr_in` can be cast to a pointer to a struct `sockaddr` and vice-versa

What's the difference between `PF_INET` and `AF_INET`???

Socket programming API: syscalls

- **int listen(int sockfd, int backlog);**
 - Put socket into passive state (wait for connections rather than initiating a connection)
 - returns 0 on success, -1 and sets errno on failure
 - **sockfd**: socket file descriptor returned by socket()
 - **backlog**: the maximum number of connections this program can serve simultaneously

Socket programming API: syscalls

- **int accept(int sockfd, struct sockaddr* client_addr, int* addrlen);**
 - Accept a new connection
 - Return client's socket file descriptor or -1. Also sets errno on failure
 - **sockfd**: socket file descriptor for server, returned by socket()
 - **client_addr**: IP address and port number of a client (returned from call)
 - **addrlen**: length of address structure = pointer to **int** set to **sizeof(struct sockaddr_in)**
 - **NOTE: client_addr and addrlen are result arguments**
 - i.e. The program passes empty client_addr and addrlen into the function, and the kernel will fill in these arguments with client's information (**why do we need them?**)

More Information about Accept()

- A new socket is cloned from the listening socket
- If there are no incoming connection to accept
 - **Non-Blocking mode:** accept() returns -1 and throw away the new socket
 - **Blocking mode (default):** accept operation was added to the wait queue

Socket programming API: syscalls

- **int connect (int sockfd, struct sockaddr* server_addr, int addrlen);**
 - Connector to another socket (server)
 - Return 0 on success, -1 and sets errno on failure
 - **sockfd**: socket file descriptor (returned from socket)
 - **server_addr**: IP address and port number of the server
 - server's IP address and port number should be known in advance
 - **addrlen**: sizeof(struct sockaddr_in)

Socket programming API: syscalls

- **int write(int sockfd, char* buf, size_t nbytes);**
 - Write data to a TCP stream
 - Return the number of sent bytes or -1 on failures
 - **sockfd**: socket file descriptor from socket ()
 - **buf**: data buffer
 - **nbytes**: the number of bytes that caller wants to send

Socket programming API: syscalls

- **int read(int sockfd, char* buf, size_t nbytes);**
 - Read data from TCP stream
 - Return the number of bytes read or -1 on failures
 - Return 0 if socket is closed
 - **sockfd**: socket file descriptor returned from socket ()
 - **buf**: data buffer
 - **nbytes**: the number of bytes that caller can read (usually set as buffer size)

Socket programming API: syscalls

- **int close(int sockfd);**
 - close a socket
 - return 0 on success, or -1 on failure
 - After close, sockfd is no longer valid

How to write a server: headers

```
/* PLEASE include these headers */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <errno.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <sys/wait.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <unistd.h>

#define MYPOR 5000 /* Avoid reserved ports */
#define BACKLOG 10 /* pending connections queue size */
```

How to write a server: body (I)

```
int main()
{
    int sockfd, new_fd; /* listen on sockfd, new connection on new_fd */
    struct sockaddr_in my_addr; /* my address */
    struct sockaddr_in their_addr; /* connector addr */
    socklen_t sin_size;

    /* create a socket */
    if ((sockfd = socket(PF_INET, SOCK_STREAM, 0)) == -1) {
        perror("socket");
        exit(1);
    }
}
```

How to write a server: body (II)

```
// ...
/* set the address info */
my_addr.sin_family = AF_INET;
my_addr.sin_port = htons(MYPORT); /* short, network byte order */
my_addr.sin_addr.s_addr = htonl(INADDR_ANY);
/* INADDR_ANY allows clients to connect to any one of the host's IP
address. Optionally, use this line if you know the IP to use:
    my_addr.sin_addr.s_addr = inet_addr("127.0.0.1");
*/
memset(my_addr.sin_zero, '\0', sizeof(my_addr.sin_zero));

/* bind the socket */
if (bind(sockfd, (struct sockaddr *) &my_addr,
        sizeof(struct sockaddr)) == -1) {
    perror("bind");
    exit(1);
}
```

How to write a server: body (III)

```
// ...
    if (listen(sockfd, BACKLOG) == -1) {
        perror("listen");
        exit(1);
    }

    while (1) { /* main accept() loop */
        sin_size = sizeof(struct sockaddr_in);
        if ((new_fd = accept(sockfd, (struct sockaddr*)
                             &their_addr, &sin_size)) == -1) {
            perror("accept");
            continue;
        }
        printf("server: got connection from %s\n",
              inet_ntoa(their_addr.sin_addr));
        close(new_fd);
    }
}
```

How to write a client?

```
/* include all the headers */
int main() {
    int sockfd, new_fd; /* listen on sockfd, new connection on new_fd */
    struct sockaddr_in my_addr; /* my address */
    struct sockaddr_in their_addr; /* connector addr */
    struct hostent* he;
    int sin_size;

    if ((sockfd = socket(PF_INET, SOCK_STREAM, 0)) == -1) {
        perror ("socket");
        exit (1);
    }

    their_addr.sin_family = AF_INET; /* interp'd by host */
    their_addr.sin_port = htons (PORT);
    their_addr.sin_addr = *((struct in_addr*) he->h_addr);
    memset(their_addr.sin_zero, '\0', sizeof their_addr.sin_zero);

    if(connect(sockfd, (struct sockaddr*) &their_addr, sizeof(struct sockaddr)) == -1) {
        perror ("connect");
        exit (1);
    }
    return 0;
}
```


Summary: what we have learned today

- What is the model for network programming?
 - **Client-Server model**
- Where are we programming?
 - **TCP and UDP in a nutshell**
- Which APIs can we use? How to use them?
 - **Socket programming**

Further Reading

- Stevens, W. Richard, Bill Fenner, and Andrew M. Rudoff. *UNIX Network Programming: The Sockets Networking API*. Vol. 1. Addison-Wesley Professional, 2004.
- Beej's Guide to Network Programming (<http://beej.us/guide/bgnet>)
- Socket Programming from Dartmouth,
<http://www.cs.dartmouth.edu/~campbell/cs60/socketprogramming.html>
- C/C++ reference: <http://en.cppreference.com>

Q&A on lectures