

# CS118 Discussion 1B, Week 7

---

Boyan Ding

# Outline

---

- Network data plane
  - NAT, IPv6
- Network control plane
  - Routing
    - Link state routing
    - Distance vector routing


# NAT (network address translation)

---

- Depletion of IPv4 addresses — short-term solution
- Use private IP addresses
- Side-benefit: security
- How to achieve?
  - <public IP:port> — <private IP:port> mapping

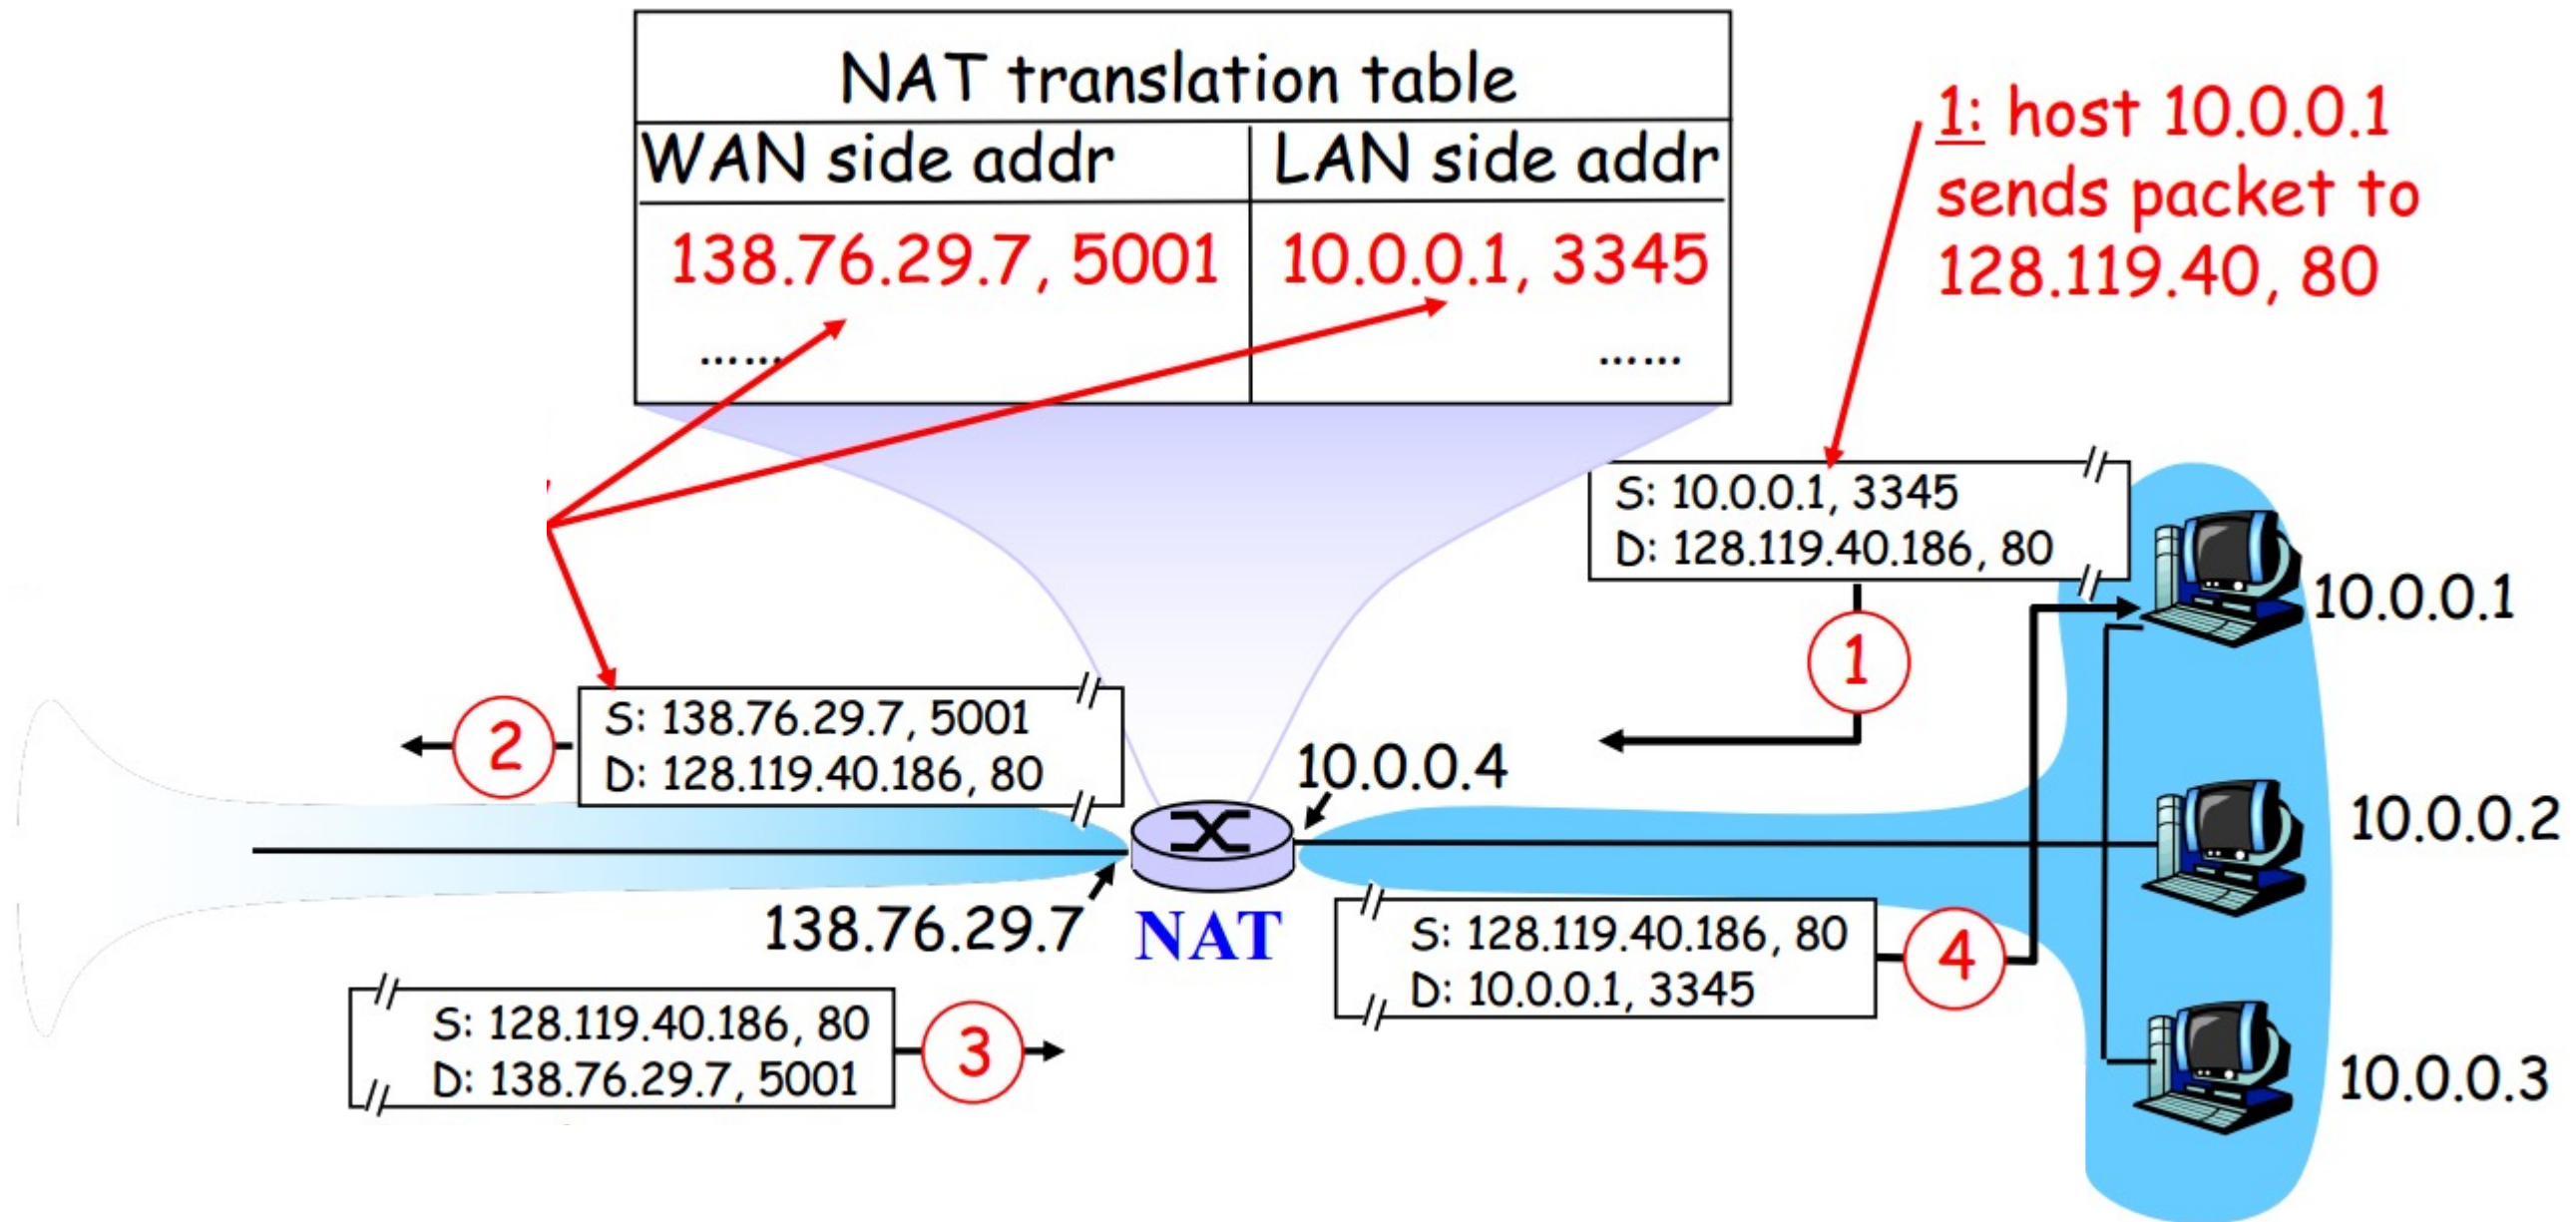
# Quick question

- 



# Quick question

- 



# NAT: Questions

---

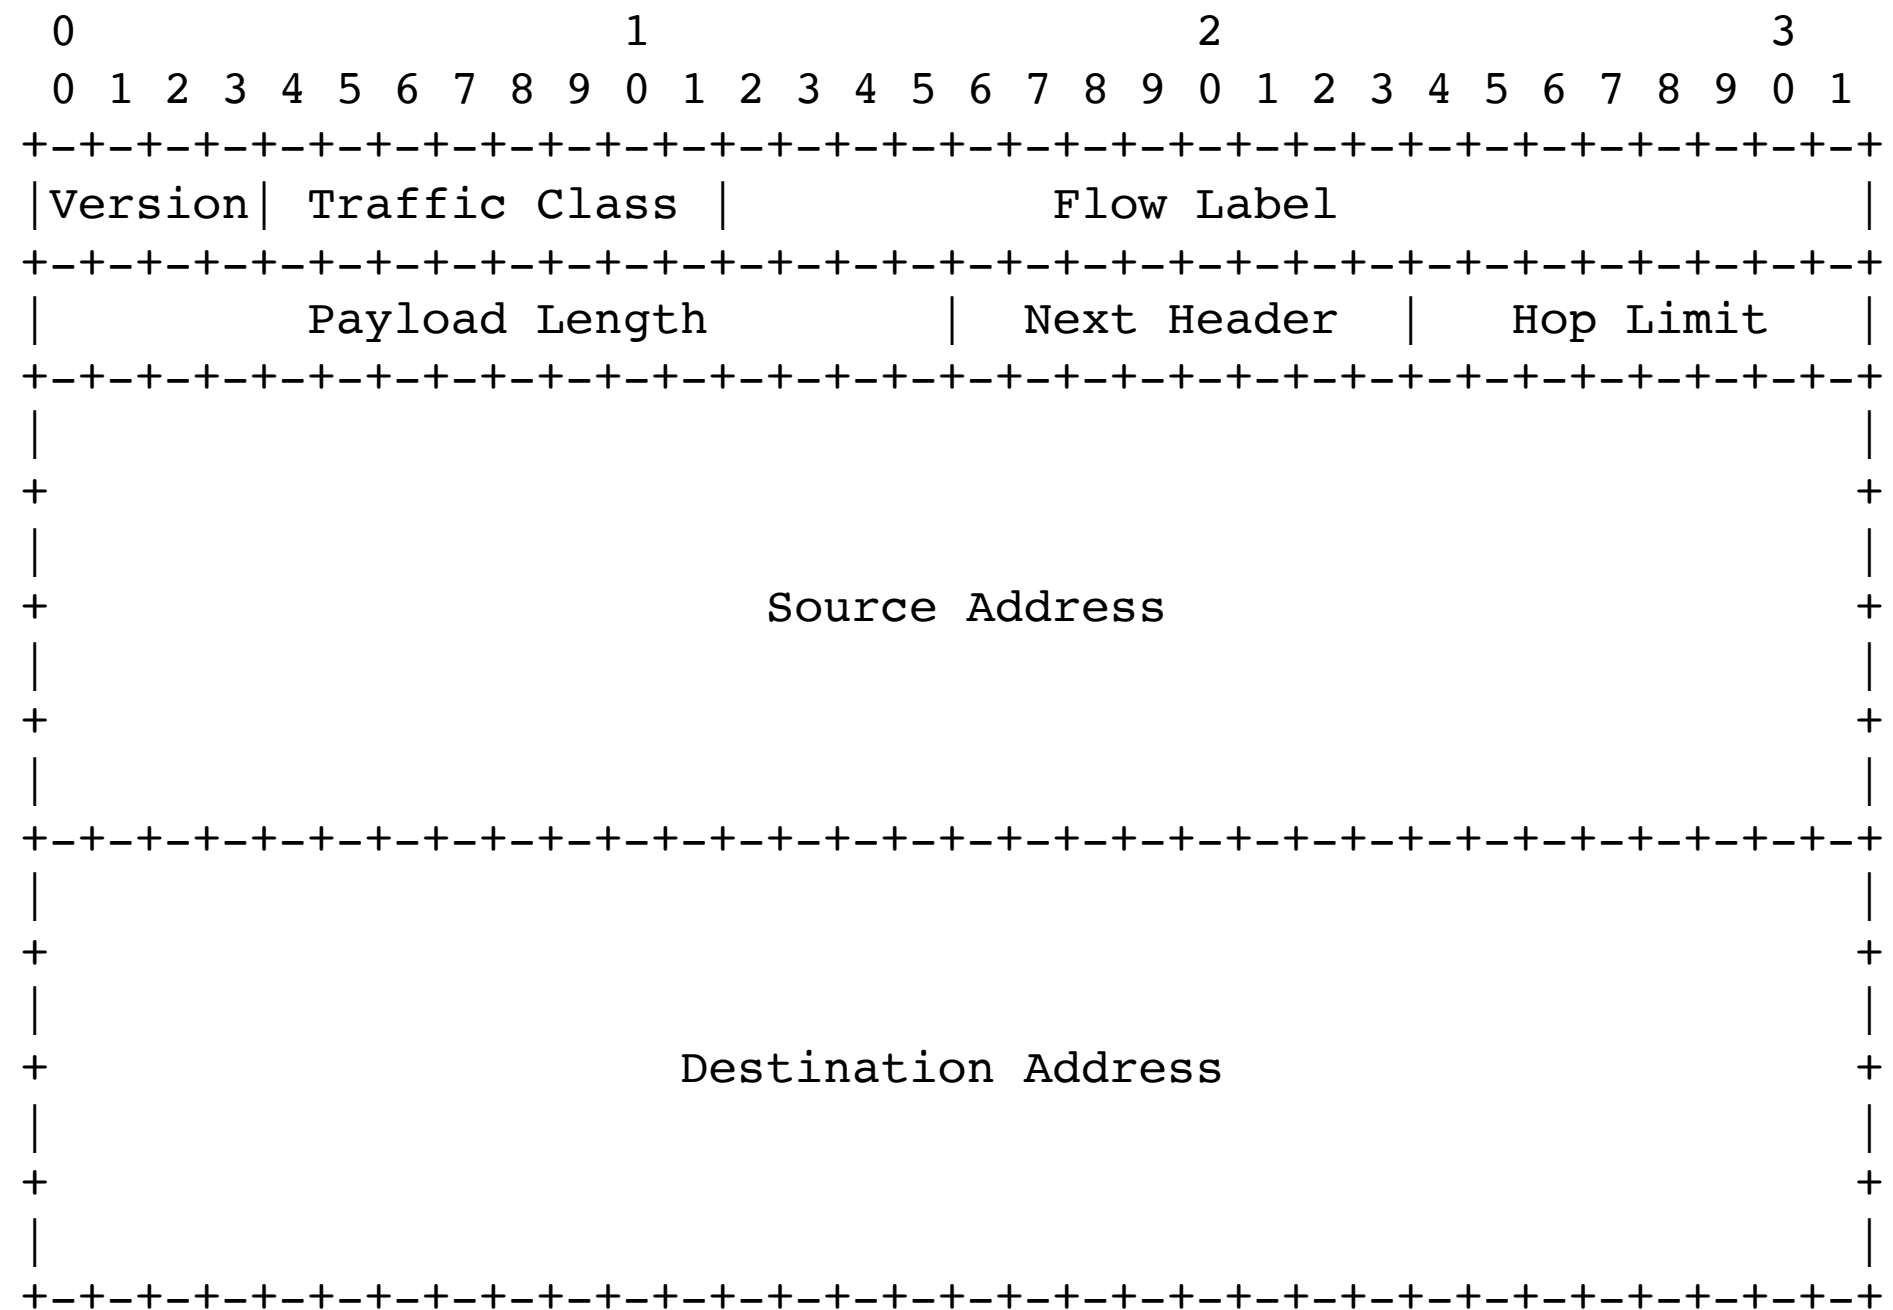
- What is the difference between an Internet router and a NAT “router” (we also call it NAT box or NAT gateway)?
  - Function? Behavior? Protocol?
- Port translation is important in NAT for TCP and UDP to work.
- Does ICMP (another transport protocol on IP mainly for network debugging e.g. ping/traceroute), which doesn't have port notation work with NAT? How?

# NAT: downside

---

- Increased complexity
- Cannot easily extend to new transport protocols
- Single point of failure
- Can hardly run services inside a NAT box
  - Why?

# IPv6



IPv6 Header Format (RFC 2460)



# IPv6/IPv4 differences

---

- Fixed-length 40 byte header
  - length field excludes header
  - Header Length field eliminated
- Address length: 128 bits
- Priority: usage yet to be finalized
- Flow Label: identify packets in same flow
- Next header: identify upper layer protocol for data
- Options: outside of the basic header, indicated by Next Header field
- Header Checksum: removed

# IPv6 address format (optional)

---

- Colon-Hex: 2607:F010:03f9:0000:0000:0000:0004:0001
- Can skip leading zeros of each word:  
2607:F010:3f9:0:0:0:4:1
- Can skip one sequence of zero words (compressed representation), e.g., 2607:f010:3f9::4:1
- Can leave the last 32 bits in dot-decimal:  
2607:f010:3f9::0.4.0.1
- Can specify a prefix by /length: 2607:f010:3f9::/64

# Special IPv6 addresses (optional)

---

- `::/128` - Unspecified
- `::1/128` - Loopback
- `::ffff:0:0/96` - IP4-mapped address
- `2002::/16` - 6to4
- `ff00::/8` - Multicast
- `fe80::/10` - Link-Local Unicast

# Routing: concepts

---

- What is the purpose of routing algorithms?
- Why do routing algorithms use shortest path?
  - What does the “distance” mean in reality?

# Routing: concepts

---

- Global or decentralized information?
  - global: all routers have complete topology, link cost info
    - algorithm?

# Routing: concepts

---

- Global or decentralized information?
  - global: all routers have complete topology, link cost info
    - “link state” algorithms

# Routing: concepts

---

- Global or decentralized information?
  - global: all routers have complete topology, link cost info
    - “link state” algorithms
  - decentralized: router knows physically-connected neighbors, link costs to neighbors; iterative process of computation, exchange of info with neighbors
    - algorithm?

# Routing: concepts

---

- Global or decentralized information?
  - global: all routers have complete topology, link cost info
    - “link state” algorithms
  - decentralized: router knows physically-connected neighbors, link costs to neighbors; iterative process of computation, exchange of info with neighbors
    - “distance vector” algorithms



# Link state routing

---

- Dijkstra's algorithm
  - net topology, link costs known to all nodes
  - computes least cost paths from one node ('source') to all other nodes
  - iterative: after  $k$  iterations, know least cost path to  $k$  destinations

# Link state routing: algorithm

---

```
1  Initialization:
2    N' = {u}
3    for all nodes v
4      if v adjacent to u
5        then D(v) = c(u,v)
6      else D(v) = ∞
7
8  Loop
9    find w not in N' such that D(w) is a minimum
10   add w to N'
11   update D(v) for all v adjacent to w and not in N':
12   [Link cost update heuristic from Dijkstra algo.]
13  until all nodes in N'
```

$c(x, y)$ : link cost from node  $x$  to  $y$ ;  $c(x, y) = \infty$  if not direct neighbors

$D(v)$ : current value of cost of path from source to destination  $v$

$p(v)$ : predecessor node along path from source to  $v$

$N'$ : set of nodes whose least cost path definitively known

# Link state routing: algorithm

---

```
1  Initialization:
2      N' = {u}
3      for all nodes v
4          if v adjacent to u
5              then D(v) = c(u,v)
6              else D(v) = ∞
7
8  Loop
9      find w not in N' such that D(w) is a minimum
10     add w to N'
11     update D(v) for all v adjacent to w and not in N':
12         D(v) = min( D(v), D(w) + c(w,v) )
13 until all nodes in N'
```

$c(x, y)$ : link cost from node  $x$  to  $y$ ;  $c(x, y) = \infty$  if not direct neighbors

$D(v)$ : current value of cost of path from source to destination  $v$


$p(v)$ : predecessor node along path from source to  $v$

$N'$ : set of nodes whose least cost path definitively known

# Link state routing: example

---

- Using link state routing to setup a forwarding table for node **u**



# Let's work it out

N'	D(v), p(v)	D(w), p(w)	D(x), p(x)	D(y), p(y)	D(z), p(z)
u	2, u	5, u	1, u	$\infty$	$\infty$
ux	2, u	4, x		2, x	$\infty$
uxy	2, u	3, y			4, y
uxyv		3, y			4, y
uxyvw					4, y
uxyvwz					

# Let's work it out

N'	D(v), p(v)	D(w), p(w)	D(x), p(x)	D(y), p(y)	D(z), p(z)
u	2, u	5, u	1, u	$\infty$	$\infty$
ux	2, u	4, x		2, x	$\infty$
uxy	2, u	3, y			4, y
uxyv		3, y			4, y
uxyvw					4, y
uxyvwz					

# Link state routing: complexity

---

- size:  $n$  nodes
- each iteration: need to check all nodes,  $w$ , not in  $N$
- $n(n+1)/2$  comparisons:  $O(n^2)$
- more efficient implementations possible:  $O(n \log n)$

# Distance vector routing

---

- Bellman-Ford equation (dynamic programming)
- let
- $dx(y) := \text{cost of least-cost path from } x \text{ to } y$
- then
- $dx(y) = ?$



# Distance vector routing


---

- Bellman-Ford equation (dynamic programming)
- let
- $dx(y) := \text{cost of least-cost path from } x \text{ to } y$
- then
- $dx(y) = \min_v \{c(x,v) + dv(y)\}, v: \text{neighbors of } x$

# Distance vector routing: example

---


- What's the cost of least-cost path for  $u \rightarrow z$ ?



# Let's work it out

---


- clearly:
  - $dv(z) = ?$ ,  $dx(z) = ?$ ,  $dw(z) = ?$



# Let's work it out

---

- clearly:
  - $dv(z) = 5$ ,  $dx(z) = 3$ ,  $dw(z) = 3$
- According to B-F equation:
  - $du(z) = \min \{ ? \}$



# Let's work it out

---

- clearly:
  - $dv(z) = 5, dx(z) = 3, dw(z) = 3$
- According to B-F equation:
  - $du(z) = \min \{c(v, x) + dv(z), c(u, x) + dx(z), c(u, w) + dw(z)\}$

# Let's work it out

---

- clearly:
  - $dv(z) = 5, dx(z) = 3, dw(z) = 3$
- According to B-F equation:
  - $du(z) = \min \{c(u, v) + dv(z), c(u, x) + dx(z), c(u, w) + dw(z)\}$ 
    - $= \min \{2 + 5, 1 + 3, 5 + 3\} = 4$

# Distance vector routing: key idea


---

- from time-to-time, each node sends its own distance vector estimate to neighbors
- when x receives new DV estimate from neighbor, it updates its own DV using B-F equation.

# Distance vector routing: caveat

---


- Count-to-infinity problem.
- Can you work out an example?






# Distance vector routing: caveat

- Count-to-infinity problem.
- Can you work out an example?



# Distance vector routing: caveat


- Count-to-infinity problem.
- Can you work out an example?



# Distance vector routing: caveat

- Count-to-infinity problem.
- Can you work out an example?
- Can you propose a solution?
- basic idea?

**A should not propagate its distance to B!**



# Distance vector routing: split horizon


---

- Previous solution idea:
  - split horizon
    - if A reaches C through B, A should not tell B that B can reach C
    - Then B will not attempt to go through A to reach C
  - Are we good?

# Distance vector routing: split horizon

---

- Previous solution idea:
  - split horizon
    - if A reaches C through B, A should not tell B that A can reach C
    - Then B will not attempt to go through A to reach C
  - Are we good?



# Distance vector routing: poison reverse



---

- Split horizon + poison reverse
  - if A reaches D through C:
    - A tells C that A's distance to D is infinite
    - Then C will not attempt to go through A to reach D
    - In practice, infinite == 16 hops (RIP protocol)

# Distance vector routing: poison reverse

If Z routes through Y to get to X:

- Z tells Y its (Z's) distance to X is infinite (so Y won't route to X via Z)



# Link State v.s. Distance Vector

	Link state	Distance vector
message complexity	with $n$ nodes, $E$ links, $O(nE)$ msgs sent	exchange between neighbors only (convergence time varies)
convergence speed	$O(n^2)$ algorithm requires $O(nE)$ msgs	convergence time varies (may be routing loops)
robustness	node can advertise incorrect link cost; each node computes only its own table	DV node can advertise incorrect path cost; error propagate thru network
implementation	OSPF	RIP



# Summary

---

- Link-state routing (Dijkstra) algorithm:
  - each node computes the shortest paths to all the other nodes based on the complete topology map
- Distance Vector (Bellman-Ford) routing algorithm:
  - each node computes the shortest paths to all the other nodes based on its neighbors distance to all destinations